

Support Vector Machines (SVM)

Jérémie Gince

Département de physique
Université de Sherbrooke

September 12, 2023

Contents

1 Binary classification

2 SVM

3 SVM in Python

4 SVM - The origin

Context




Figure: Two classes being separate by a line with the parameters W and b and the estimator $h(\cdot)$.

Hinge Loss

The hinge loss used to maximize the margins of a binary classifier is the following:

$$E_{\text{hinge}}(\theta|\mathcal{D}) = \sum_{\{x^t, y^t\} \in \mathcal{D}} \max(0, 1 - y^t h(x^t|\theta)) + E_{\text{regularization}}(\theta). \quad (1)$$

- if $y^t = 1$ and $h^t > 0 \rightarrow$ good classification with $E \in [0, 1)$
- if $y^t = 1$ and $h^t < 0 \rightarrow$ bad classification with $E > 0$
- if $y^t = -1$ and $h^t < 0 \rightarrow$ good classification with $E \in [0, 1)$
- if $y^t = -1$ and $h^t > 0 \rightarrow$ bad classification with $E > 0$

Estimator

The linear separator is defined by the decision function

$$h^t(x^t|W, b) = \sum_{j=1}^M W_j x_j^t + b \quad (2)$$

$$\hat{y}^t = \phi(h^t(x^t|W, b)) \quad (3)$$

with a training dataset $\mathcal{D} = \{x^t, y^t\}$ where

- x_j^t is j -th entry of the input t of size M ;
- W_j is the j -th weight from the vector W of size M ;
- b is the bias as a scalar;
- $\phi(\cdot)$ is the activation function which is often $\text{sign}(\cdot)$ in a binary classifier;
- y^t is the target class t ;
- \hat{y}^t is the predicted class t .

Changing the decision space

It's possible that the data \mathcal{D} are not linearly separable in the initial space \mathcal{S}_0 but they are in another space \mathcal{S}_1 . Then we can use a transformation $\psi : \mathbb{R}^M \mapsto \mathbb{R}^N$.

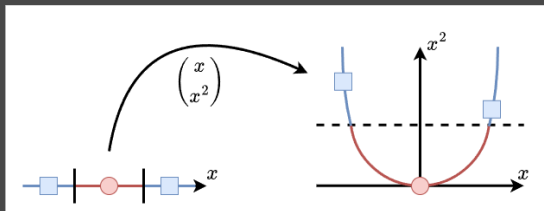


Figure: Two classes not separable in initial space but are separable after the transformation $\psi(x) = (x \ x^2)^T$.

- The linear separator will be

$$h(x) = W^T \psi(x) + b.$$

SVM

In a SVM, the parameters W of the separator are

$$W = \sum_t \alpha^t y^t \psi(x^t) \quad (5)$$

with α^t , a parameter associated with x^t . This lead to

$$h(x) = \sum_t \alpha^t y^t (\psi(x^t))^T \psi(x) + b. \quad (6)$$

The Kernel

The kernel function is defined as

$$K(x, y) = (\psi(x))^T \psi(y). \quad (7)$$

Using the kernel definition, the SVM decision function is

$$h(x) = \sum_t \alpha^t y^t K(x^t, x) + b. \quad (8)$$

The computing of $\psi(\cdot)$ is now hidden in the new kernel function, so this $\psi(\cdot)$ can be "unknown" from the SVM decision function.

Considering that some α^t are null, we can write the prediction function as

$$h(x) = \sum_s \alpha^s y^s K(x^s, x) + b, \quad (9)$$

$$\hat{y} = \phi(h(x)) \quad (10)$$

where

- the index s are associated with the non-zero α^t . The x^s are what we called, the support vectors.




Figure: Binary classification. The dots lines represent the margins of the classifier and the circled data are the support vectors.

Finding the parameters

To Learn the parameters α^t and b of the SVM classifier, we can do a gradient descent in L iterations. The updates are

- $\Delta\alpha^t = -\eta \frac{\partial E(\alpha, b | \mathcal{D})}{\partial \alpha^t},$
- $\Delta b = -\eta \frac{\partial E(\alpha, b | \mathcal{D})}{\partial b}$

with the user-defined error function $E(\cdot)$ and the learning rate η . The update of the parameters will be

$$\alpha_\kappa = \begin{cases} 0 & \text{if } \alpha_{\kappa-1} + \Delta\alpha^t < 0 \\ \alpha_{\kappa-1} + \Delta\alpha^t & \text{else,} \end{cases} \quad (11)$$

$$b_\kappa = b_{\kappa-1} + \Delta b, \quad \forall \kappa \in \{1, \dots, L\}. \quad (12)$$

Hinge Loss

The hinge loss is generally used for SVM.

$$E_{\text{hinge}}(\alpha, b|\mathcal{D}) = \sum_{\{x^t, y^t\} \in \mathcal{D}} \max(0, 1 - y^t h(x^t|\alpha, b)) + \lambda \frac{1}{2} \sum_{\alpha^s \in \alpha} (\alpha^s)^2 \quad (13)$$

Code

- Repository: https://github.com/JeremieGince/Learning_SVM
- Google Colab: [Learning_SVM.google.colab](https://colab.research.google.com/notebooks/Learning_SVM.ipynb)

Support vector machines (SVM)

SVM is a special type of linear separator with following goals:

- Maximize the distance between the separating hyperplane and the training data:

$$d_{h^t, x^t} = \frac{|W^T x^t + b|}{\|W\|} = \frac{y^t (W^T x^t + b)}{\|W\|} \quad (14)$$

- Find the weights W and b in order to have

$$d_{h^t, x^t} \geq \rho, \forall t \quad (15)$$

where ρ is the margin.

Lagrange multipliers

In order to achieve those goals, we have to use the Lagrange multipliers.

- It's a method used to optimize a function $f(x)$ under a constraint $g(x)$ written in a way that $g(x) = 0$,
- with the Lagrangian

$$\mathcal{L}(x, \alpha) = f(x) + \alpha g(x). \quad (16)$$

- The stationary points of the Lagrangian and the optimal solution of $f(x)$ is then found with

$$\nabla \mathcal{L}(x, \alpha) = 0. \quad (17)$$

Lagrangian multipliers - Example

Suppose we want to maximize the function $f(q_i) = q_1 + q_2$ with the constraint $q_1^2 + q_2^2 = 1$. The constraint will then be noted as

$$g(q_i) = q_1^2 + q_2^2 - 1 = 0, \quad (18)$$

and the Lagrangian as

$$\mathcal{L}(q_i, \alpha) = f(q_i) + \alpha g(q_i) \quad (19)$$

$$\implies \mathcal{L}(q_i, \alpha) = q_1 + q_2 + \alpha(q_1^2 + q_2^2 - 1). \quad (20)$$

The gradient of the Lagrangian will be

$$\nabla \mathcal{L}(q_i, \alpha) = \left[\frac{\partial \mathcal{L}(q_i)}{\partial q_1}, \quad \frac{\partial \mathcal{L}(q_i)}{\partial q_2}, \quad \frac{\partial \mathcal{L}(q_i)}{\partial \alpha} \right], \quad (21)$$

$$= [1 + 2\alpha q_1, \quad 1 + 2\alpha q_2, \quad q_1^2 + q_2^2 - 1], \quad (22)$$

$$= \begin{bmatrix} 1 + 2\alpha q_1 = 0, \\ 1 + 2\alpha q_2 = 0, \\ q_1^2 + q_2^2 - 1 = 0 \end{bmatrix}, \quad (23)$$

$$q_1 = q_2 = -\frac{1}{2\alpha}, \quad (24)$$

$$\Rightarrow \alpha^{(1)} = +\frac{1}{\sqrt{2}}, \quad \alpha^{(2)} = -\frac{1}{\sqrt{2}}. \quad (25)$$

Since $f(-\frac{1}{2\alpha^{(1)}}, -\frac{1}{2\alpha^{(1)}}) = -\sqrt{2}$ and $f(-\frac{1}{2\alpha^{(2)}}, -\frac{1}{2\alpha^{(2)}}) = \sqrt{2}$, the Lagrange multiplier that maximize our function $f(q_i)$ is $\alpha^{(2)}$.

Lagrange multipliers with inequality

If the constraint is $g(x) \geq 0$, the conditions for optimality will be

- $g(x) \geq 0$;
- $\alpha \geq 0$;
- $\alpha g(x) = 0$.

Then to minimize $f(x)$ with $g(x) \geq 0$, we have to optimize the Lagrangian

$$\mathcal{L}(x, \alpha) = f(x) - \alpha g(x) \quad (26)$$

with $\alpha \geq 0$.

SVM

The problem:

- Minimize $\frac{1}{2} \|W\|^2$
- with the constraint $y^t (W^T x^t + b) \geq 1 \ \forall t$

Re-writing using the Lagrangian multiplier:

$$\mathcal{L}_p = \frac{1}{2} \|W\|^2 - \sum_t \alpha^t [y^t (W^T x^t + b) - 1] \quad (27)$$

$$= \frac{1}{2} \|W\|^2 - \sum_t \alpha^t y^t (W^T x^t + b) + \sum_t \alpha^t \quad (28)$$

with \mathcal{L}_p , the primal form of the Lagrangian.

Simplification - dual form

Since we want to find an stationary point of \mathcal{L}_p :

$$\frac{\partial \mathcal{L}_p}{\partial W} = 0 = W - \sum_t \alpha^t y^t x^t, \quad (29)$$

$$\implies W = \sum_t \alpha^t y^t x^t, \quad (30)$$

and

$$\frac{\partial \mathcal{L}_p}{\partial b} = 0 = \sum_t \alpha^t y^t. \quad (31)$$

The dual form can be written as

$$\mathcal{L}_d = \frac{1}{2}(W^T W) - \underbrace{W^T \sum_t \alpha^t y^t x^t}_W - \underbrace{b \sum_t \alpha^t y^t}_0 + \sum_t \alpha^t \quad (32)$$

$$= -\frac{1}{2}(W^T W) + \sum_t \alpha^t \quad (33)$$

$$\implies \mathcal{L}_d = -\frac{1}{2} \sum_t \sum_s \alpha^t \alpha^s y^t y^s (x^t)^T x^s + \sum_t \alpha^t \quad (34)$$

Dual form

Re-writing the problem using the dual form.

- Maximize $-\frac{1}{2} \sum_t \sum_s \alpha^t \alpha^s y^t y^s (x^t)^T x^s + \sum_t \alpha^t$
- With the constraints $\sum_t \alpha^t y^t = 0$ and $\alpha^t \geq 0 \forall t$
- In practice, a majority of α^t will be zero;
- The data x^t with $\alpha^t > 0$ are called the support vectors.

Changing the decision space

It's possible that the data \mathcal{D} are not linearly separable in the initial space \mathcal{S}_0 but are in another space \mathcal{S}_1 . Then we can use a transformation $\psi : \mathbb{R}^M \mapsto \mathbb{R}^N$ that $x \mapsto \psi(x)$.

- The linear separator will be

$$h(x) = W^T \psi(x) + b \quad (35)$$

- In the dual form:

$$W = \sum_t \alpha^t y^t \psi(x^t) \quad (36)$$

$$h(x) = \sum_t W^T \psi(x) + b = \sum_t \alpha^t y^t (\psi(x^t))^T \psi(x) + b \quad (37)$$

The Kernel

The kernel function is defined as

$$K(x, y) = (\psi(x))^T \psi(y). \quad (38)$$

Using the kernel definition, the SVM decision function is

$$h(x) = \sum_t \alpha^t y^t K(x^t, x) + b. \quad (39)$$

The computing of $\psi(\cdot)$ is now hidden in the new kernel function, so this $\psi(\cdot)$ can be "unknown" from the SVM decision function.

The resulting prediction function of the SVM:

$$h(x) = \sum_{x^s} \alpha^s y^s K(x^s, x) + b, \quad (40)$$

$$\hat{y} = \phi(h(x)). \quad (41)$$

- Reminder: x^s is the support vector s associated with the non-zero Lagrangian multiplier α^s .

The End